



Troubleshooting Pentaho Business Analytics: Common Problems and Solutions



This document is copyright © 2012 Pentaho Corporation. No part may be reprinted without written permission from Pentaho Corporation. All trademarks are the property of their respective owners.

Help and Support Resources

If you have questions that are not covered in this guide, or if you would like to report errors in the documentation, please contact your Pentaho technical support representative.

Support-related questions should be submitted through the Pentaho Customer Support Portal at <http://support.pentaho.com>.

For information about how to purchase support or enable an additional named support contact, please contact your sales representative, or send an email to sales@pentaho.com.

For information about instructor-led training on the topics covered in this guide, visit <http://www.pentaho.com/training>.

Limits of Liability and Disclaimer of Warranty

The author(s) of this document have used their best efforts in preparing the content and the programs contained in it. These efforts include the development, research, and testing of the theories and programs to determine their effectiveness. The author and publisher make no warranty of any kind, express or implied, with regard to these programs or the documentation contained in this book.

The author(s) and Pentaho shall not be liable in the event of incidental or consequential damages in connection with, or arising out of, the furnishing, performance, or use of the programs, associated instructions, and/or claims.

Trademarks

Pentaho (TM) and the Pentaho logo are registered trademarks of Pentaho Corporation. All other trademarks are the property of their respective owners. Trademarked names may appear throughout this document. Rather than list the names and entities that own the trademarks or insert a trademark symbol with each mention of the trademarked name, Pentaho states that it is using the names for editorial purposes only and to the benefit of the trademark owner, with no intention of infringing upon that trademark.

Company Information

Pentaho Corporation
Citadel International, Suite 340
5950 Hazeltine National Drive
Orlando, FL 32822
Phone: +1 407 812-OPEN (6736)
Fax: +1 407 517-4575
<http://www.pentaho.com>

E-mail: communityconnection@pentaho.com

Sales Inquiries: sales@pentaho.com

Documentation Suggestions: documentation@pentaho.com

Sign-up for our newsletter: <http://community.pentaho.com/newsletter/>

Contents

| | |
|--|----|
| Introduction..... | 5 |
| General..... | 6 |
| File Names and Paths..... | 6 |
| JDBC Driver Problems..... | 6 |
| Adding a JDBC Driver..... | 6 |
| Data Conversion Issues With MySQL Driver 5.0.8..... | 7 |
| Fixing JTDS varchar(MAX) Limitations in MSSQL 2005..... | 7 |
| Windows Domains Won't Authenticate When Using the JTDS Driver..... | 7 |
| Version Check..... | 7 |
| I don't know what the default login is for the DI Server, Enterprise Console, and/or Carte..... | 8 |
| Examining Log Files..... | 8 |
| Upgrade..... | 9 |
| Tomcat Logs Report Memory Leaks..... | 9 |
| context.xml Changes Do Not Take Effect After Re-deploying a WAR..... | 9 |
| javax.jcr.RepositoryException: no search manager configured for this workspace..... | 9 |
| User Console Themes Render Improperly After Upgrade..... | 10 |
| BA Server and Pentaho Enterprise Console..... | 11 |
| Address Already In Use: JVM_Bind When Starting Enterprise Console..... | 11 |
| Library Conflicts..... | 11 |
| Report Parameters That Include Accented Characters Fail to Validate..... | 11 |
| Modifying server.xml To Work With Accented Characters..... | 11 |
| vfs-provider.xml Duplicates..... | 12 |
| Varying Context and Data Source Configuration Methods..... | 12 |
| Licenses Not Found After Installation..... | 12 |
| Evaluation Licenses are Not Granted..... | 12 |
| Cannot Create Hibernate Tables in MySQL..... | 12 |
| Case-Insensitivity for Usernames, Passwords, and Filenames..... | 13 |
| Unable to Use the Database Init Scripts for PostgreSQL..... | 13 |
| HTTP 500 or "Unable to connect to BA Server" Errors When Trying to Access Enterprise Console..... | 13 |
| Configuring the Proxy Trusting Filter..... | 13 |
| JBoss Fails to Start When the Pentaho HSQLDB Sample Database Is Running..... | 14 |
| JBoss Fails to Start After Manually Unpacking pentaho.war..... | 14 |
| Web-App Path doesn't validate in Enterprise Console..... | 14 |
| LDAP incorrectly authenticates user IDs that don't match letter case..... | 14 |
| Forcing Case-Sensitivity in LDAP..... | 14 |
| LDAP Roles Are Not "Admin" and "Authenticated"..... | 15 |
| With LDAP Authentication, the PDI Repository Explorer is Empty..... | 15 |
| Old Action Sequences Fail on Scrollable Result Sets..... | 16 |
| Geo Maps partially rendered..... | 16 |
| Setting the Geo Map Rendering Wait Time..... | 16 |
| Mac OSX Error When Running Apps..... | 16 |
| Report Designer and Reporting Engine..... | 18 |
| Enabling Multi-Value Report Parameters for Metadata-based Queries Created with Previous Versions of Report Designer..... | 18 |
| Report Elements With Dynamic Heights Overlap Other Elements..... | 18 |
| Null Pointer Exception When Sharing Result Sets..... | 18 |
| Columns Unexpectedly Merge When Exporting to Excel..... | 18 |
| Analysis..... | 19 |
| Old Analysis Schemas Still Show Up in Pentaho User Console..... | 19 |
| Removing Mondrian Data Sources..... | 19 |
| Multi-Byte Characters Don't Appear In PDFs Exported From Analyzer..... | 19 |
| Setting a Default Font for PDF Exports..... | 20 |
| Data Integration..... | 21 |
| Troubleshooting a Corrupted DI Server Repository..... | 21 |
| Using the H2 Database Recovery Tool..... | 21 |
| Action Sequences That Call PDI Content Won't Run..... | 22 |

| | |
|---|----|
| Adding PDI Enterprise Repository Content Support to the BA Server..... | 22 |
| Jobs scheduled on the DI Server cannot execute a transformation on a remote Carte server..... | 23 |
| Executing Scheduled Jobs on a Remote Carte Server..... | 23 |
| Kitchen can't read KJBs from a Zip export..... | 24 |
| PDI Transformation Logging Doesn't Show In PEC..... | 25 |
| DI Server fails to start (IOException: Premature end of file)..... | 25 |
| Using ODBC..... | 25 |
| Metadata..... | 26 |
| Managing Multiple Outer-Joins..... | 26 |
| Using the Delay Outer Join Conditions Property..... | 27 |

Introduction

This document contains no unique content; it is a compilation of all of the troubleshooting content contained in Pentaho Enterprise Edition documentation. Each guide has its own troubleshooting section, where applicable, that contains problems and solutions that Pentaho customers and partners have encountered in the past or are anticipated to encounter in the future. The *Troubleshooting Guide* is designed to address customers, partners, consultants, and internal Pentaho employees who are already familiar with the installation, configuration, and operation of Business Analytics, and only need to consult the documentation to solve a problem.

General

This section contains problems and solutions that apply to Business Analytics in general.

File Names and Paths



Note: This is the most common installation problem.

Many of the configuration files and paths in this guide are similar, and it is easy to confuse them, which could result in modifying the wrong files or copying to the wrong locations. Double-check your file names and paths and ensure that you've copied all of the right files to all of the correct directories.

Trailing slashes are important; both their inclusion and their absence, depending on the file and parameter or element you are modifying. Follow the examples in this guide exactly unless otherwise directed.

JDBC Driver Problems

First, ensure that the correct JDBC driver JARs are installed to the correct locations, then check to make sure that there aren't conflicting driver versions. See the *Pentaho Business Analytics Administrator's Guide* for explanations of how to create JDBC Database connections and driver locations for all parts of Pentaho Business Analytics. Check with your database or driver vendor if you suspect you have having JDBC driver compatibility issues.

Adding a JDBC Driver

Before you can connect to a data source in any Pentaho server or client tool, you must first install the appropriate database driver. Your database administrator, Chief Intelligence Officer, or IT manager should be able to provide you with the proper driver JAR. If not, you can download a JDBC driver JAR file from your database vendor or driver developer's Web site. Once you have the JAR, follow the instructions below to copy it to the driver directories for all of the Business Analytics components that need to connect to this data source. See the *Compatibility Matrix: Supported Components* in any of the Installation guide for current version numbers.



Note: Microsoft SQL Server users frequently use an alternative, non-vendor-supported driver called JTDS. If you are adding an MSSQL data source, ensure that you are installing the correct driver.

Backing up old drivers

You must also ensure that there are no other versions of the same vendor's JDBC driver installed in these directories. If there are, you may have to back them up and remove them to avoid confusion and potential class loading problems. This is of particular concern when you are installing a driver JAR for a data source that is the same database type as your Pentaho solution repository. If you have any doubts as to how to proceed, contact your Pentaho support representative for guidance.

Installing JDBC drivers

Copy the driver JAR file to the following directories, depending on which servers and client tools you are using (Dashboard Designer, ad hoc reporting, and Analyzer are all part of the BA Server):



Note: For the DI Server: before copying a new JDBC driver, ensure that there is not a different version of the same JAR in the destination directory. If there is, you must remove the old JAR to avoid version conflicts.

- **BA Server:** `/pentaho/server/biserver-ee/tomcat/lib/`
- **Enterprise Console:** `/pentaho/server/enterprise-console/jdbc/`
- **Data Integration Server:** `/pentaho/server/data-integration-server/tomcat/webapps/pentaho-di/WEB-INF/lib/`
- **Data Integration client:** `/pentaho/design-tools/data-integration/libext/JDBC/`
- **Report Designer:** `/pentaho/design-tools/report-designer/lib/jdbc/`
- **Schema Workbench:** `/pentaho/design-tools/schema-workbench/drivers/`
- **Aggregation Designer:** `/pentaho/design-tools/agg-designer/drivers/`

- **Metadata Editor:** /pentaho/design-tools/metadata-editor/libext/JDBC/



Note: To establish a data source in the Pentaho Enterprise Console, you must install the driver in both the Enterprise Console and the BA Server or Data Integration Server. If you are just adding a data source through the Pentaho User Console, you do not need to install the driver to Enterprise Console.

Restarting

Once the driver JAR is in place, you must restart the server or client tool that you added it to.

Connecting to a Microsoft SQL Server using Integrated or Windows Authentication

The JDBC driver supports Type 2 integrated authentication on Windows operating systems through the **integratedSecurity** connection string property. To use integrated authentication, copy the **sqljdbc_auth.dll** file to all the directories to which you copied the JDBC files.

The **sqljdbc_auth.dll** files are installed in the following location:

```
<installation directory>\sqljdbc_<version>\<language>\auth\
```



Note: Use the **sqljdbc_auth.dll** file, in the x86 folder, if you are running a 32-bit Java Virtual Machine (JVM) even if the operating system is version x64. Use the **sqljdbc_auth.dll** file in the x64 folder, if you are running a 64-bit JVM on a x64 processor. Use the **sqljdbc_auth.dll** file in the IA64 folder, you are running a 64-bit JVM on an Itanium processor.

Data Conversion Issues With MySQL Driver 5.0.8

The MySQL JDBC driver version 5.0.8 may cause data conversion errors in the BA Server. For example, SQL statements that convert a numeric field to a string are returned as a string with version 5.0.7, but return as a byte array in 5.0.8. To solve this problem, you must replace the **mysql-connector-java-5.0.8.jar** with the **mysql-connector-java-5.0.7.jar** in your client tool or application server's **lib** directory.

Fixing JTDS varchar(MAX) Limitations in MSSQL 2005

Creating a report that uses a column of type varchar(MAX) may result in a **net.sourceforge.jtds.jdbc.ClobImpl@83cf00** error when using the JTDS SQL Server driver. This is caused by inherent limitations in older versions of the JTDS driver. Additionally, the SQL Server JDBC driver version 1.2.2828 also has issues accessing a varchar(MAX) column.

The solution is to upgrade the MSSQL 2005 JDBC driver to version 1.0.809.102 or later. To do this, download and install the <http://msdn.microsoft.com/en-us/sqlserver/aa937724> file from microsoft.com, then restart your MSSQL server.

Windows Domains Won't Authenticate When Using the JTDS Driver

If you are using a JTDS JDBC driver and you want to use a Windows domain user to authenticate to a Microsoft SQL Server, be aware that the Windows syntax will not work for specifying the domain and user. In the URL field, the domain must be appended to the end of the URL with a semicolon:

```
jdbc:jtds:sqlserver://svn-devel.example.com:1533/reportsInProgress;domain=testdomain
```

Version Check

The instructions in this guide are specific to the Pentaho BA Server Enterprise Edition version 4.8-GA. The installation process can change significantly between BA Server releases to address new features, updated requirements, and bug workarounds, so the instructions in this guide should be assumed not to work with any other BA Server version, including the open source BA Server Community Edition.

I don't know what the default login is for the DI Server, Enterprise Console, and/or Carte

For the DI Server administrator, it's username **admin** and password **secret**.

For Enterprise Console administrator, it's username **admin** and password **password**.

For Carte, it's username **cluster** and password **cluster**.

Be sure to change these to new values in your production environment.



Note: DI Server users are not the same as BI Server users.

Examining Log Files

If the BA Server fails to start or work properly, the log file you should consult is **pentaho.log** in the `/pentaho/server/biserver-ee/tomcat/bin/` directory. The contents of this file will assist you in tracking down the problem.

Upgrade

This section contains problems and solutions that apply to the Upgrade Guide for both the BA Server and Pentaho Data integration.

Tomcat Logs Report Memory Leaks

When shutting down Tomcat, you may see some SEVERE-level warnings in the log file similar to these:

```
Dec 17, 2010 10:18:19 AM org.apache.catalina.loader.WebappClassLoader
clearReferencesJdbc
SEVERE: The web application [/pentaho] registered the JDBC driver
[mondrian.olap4j.MondrianOlap4jDriver] but failed to unregister it when the web
application was stopped. To prevent a memory leak, the JDBC Driver has been forcibly
unregistered.
Dec 17, 2010 10:18:19 AM org.apache.catalina.loader.WebappClassLoader
clearReferencesThreads
SEVERE: The web application [/pentaho] appears to have started a thread named [HSQldb
Timer @49cf9f] but has failed to stop it. This is very likely to create a memory
leak.
Dec 17, 2010 10:18:19 AM org.apache.catalina.loader.WebappClassLoader
clearReferencesThreads
SEVERE: The web application [/pentaho] appears to have started a thread named [MySQL
Statement Cancellation Timer] but has failed to stop it. This is very likely to
create a memory leak.
Dec 17, 2010 10:18:19 AM org.apache.catalina.loader.WebappClassLoader
clearThreadLocalMap
SEVERE: The web application [/pentaho] created a ThreadLocal
with key of type [java.lang.InheritableThreadLocal] (value
[java.lang.InheritableThreadLocal@a1320e]) and a value of type
[org.pentaho.platform.engine.security.session.TrustedSystemStartupSession] (value
[org.pentaho.platform.engine.security.session.TrustedSystemStartupSession@111089b])
but failed to remove it when the web application was stopped. This is very likely to
create a memory leak.
```

These warnings are nothing to be concerned about when shutting down the Tomcat server, since they report problems with processes that are immanently being killed. However, they can have significance if you are only restarting or redeploying the Pentaho BA Server or DI Server Web applications. To avoid any memory leak problems in redeployment, you should restart Tomcat instead of redeploying or restarting the Web application with a live server.

context.xml Changes Do Not Take Effect After Re-deploying a WAR

Re-deployment of a WAR or EAR with a custom **context.xml** may cause the original context.xml that you deployed with the original WAR or EAR to become permanently cached. Tomcat in particular will generate a WAR-specific context configuration file, and keep it in place even after the WAR is deleted. The location and naming convention for this file are: **\$CATALINA_HOME/conf/Catalina/<host>/<war name>.xml**. Typically this will be something like: **/tomcat/conf/Catalina/localhost/pentaho.xml**. If this file exists, you will have to delete it prior to re-deploying pentaho.war if you have made any changes to context.xml.

javax.jcr.RepositoryException: no search manager configured for this workspace

If you see this error in your PDI server log, then there was an error in the upgrade process from PDI 4.2.0 to 4.3. Specifically, the **SearchIndex** XML nodes were not properly modified. To fix this, refer to the **Upgrading a Data Integration Server** piece and closely follow the instructions for modifying repository configuration files.

User Console Themes Render Improperly After Upgrade

If you are seeing strange rendering problems in the Pentaho User Console shortly after performing a BA Server upgrade, the problem may be related to old JavaScript files being held in the browser cache. To fix this problem, clear your Web browser's cache, then reload the Pentaho User Console.

BA Server and Pentaho Enterprise Console

This section contains problems and solutions that pertain to the BA Server and Pentaho Enterprise Console.

Address Already In Use: JVM_Bind When Starting Enterprise Console

If you encounter this error when starting Enterprise Console, then the default port 8088 is not available on your system. To change the Enterprise Console port, edit the `/pentaho/server/enterprise-console/resource/config/console.properties` file and change the value of **console.start.port.number** to an open port.

Library Conflicts

The BA Server relies on many third-party libraries that provide everything from database connectivity to specific Java classes that add necessary features to the BA Server. If you have incompatible versions of any of these third-party libraries in your application server's global lib directory, they can cause a variety of problems related to starting and running the BA Server. You will have to discover and individually canonicalize these files according to your needs.

Some known-problematic JARs are:

- commons-collections-3.2.jar (from Pentaho)
- commons-collections.jar (from JBoss in `/jboss/server/default/lib/`)
- jettison-1.01.jar (from Pentaho)
- jettison.jar (from JBoss in `/jboss/default/deploy/jbossaws.sar`)

Report Parameters That Include Accented Characters Fail to Validate

If you run a report that has parameters that include accented characters and you see an error message that says "This parameter value is of an invalid value," then you must make the Tomcat server modification explained in this task: [Modifying server.xml To Work With Accented Characters](#) on page 11.

Modifying server.xml To Work With Accented Characters

This procedure is only necessary if you plan to use character sets that include accents, such as Spanish and French.

Follow the instructions below to implement accented character set support. Change the paths to match your configuration.

1. Stop the Tomcat service.

```
sudo /etc/init.d/tomcat stop
```

2. Open the `/tomcat/server/conf/server.xml` file with a text editor.
3. Locate each **Connector** node (typically there are four in a default Tomcat configuration) and add a **URIEncoding="UTF-8"** parameter to it, as shown below:

```
<Connector URIEncoding="UTF-8" port="8080" protocol="HTTP/1.1"
    connectionTimeout="20000"
    redirectPort="8443" />
```

4. Save and close the file, then restart the Tomcat service.

```
sudo /etc/init.d/tomcat start
```

Tomcat is now properly configured to handle accented characters.

vfs-provider.xml Duplicates

The above-referenced configuration file may be present in a number of JARs in other applications that you've deployed to your Java application server. Having multiple instances of this file will cause classpath exceptions. You must merge the multiple files into one canonical edition in order to solve the problem.

Varying Context and Data Source Configuration Methods

The instructions in this guide direct you to edit a **context.xml** override for Tomcat and build it into the WAR. Your application server may not support this method, or may not accept these files by default. This will be especially true if you are using unsupported application servers or versions, or if you are on Linux and your distribution provider alters the default package configuration for Tomcat.

There are other methods of changing context settings, especially through other kinds of configuration files such as **server.xml** or application-specific XML files that your application server creates for each deployed WAR. If the context and data source instructions in this guide do not work for you, consult your application server's documentation to learn the preferred method of configuration, and adapt the Pentaho-supplied process to accommodate it.

Licenses Not Found After Installation

If you've successfully installed the BA Server and Pentaho Enterprise Console, then installed licenses, and are later unable to access Dashboard Designer, Pentaho Analyzer, or Pentaho Interactive Reporting through the Pentaho User Console due to missing licenses, then you must either reinstall those licenses under the user account that will always start the Pentaho Enterprise Console and BA Server services, or you must set a Java virtual machine parameter to point to a static license path. This will override the default license path, which changes depending on which user started the Pentaho Enterprise Console and BA Server services or servers. The parameter is appended to the Tomcat or JBoss service execution commands, and to the Pentaho Enterprise Console **start-pec** script. Examples are shown below:

Windows Tomcat service adjustment command

```
tomcat5 //US//pentahobiserver ++JvmOptions "-Dpentaho.installed.licenses.file=C:\Pentaho\installedLicenses.xml"
```

Linux start-pec.sh script snippet

```
"$_PENTAHO_JAVA" -Dpentaho.installed.licenses.file="/home/pentaho/.pentaho/.installedLicenses.xml" -Xmx512M -XX:PermSize=64M -XX:MaxPermSize=128M -Djava.awt.headless=true -DCONSOLE_HOME=$DIR_REL -Dlog4j.configuration=resource/config/log4j.xml -cp $CLASSPATH com.pentaho.pac.server.ProJettyServer
```

Evaluation Licenses are Not Granted

If you see a message in the graphical installation utility's error log that says, **Evaluation licenses were not granted for the following products**, this means that you've installed this version of Pentaho Business Analytics (or some piece of it) in the past, and one or more evaluation licenses have already been generated and installed. An evaluation license expires 30 days after it is generated; once expired, it cannot be reactivated. The only way to use Pentaho Business Analytics's full functionality is to install a new, active license that overwrites the old one.

In some cases, you may have delayed or extended your Pentaho evaluation beyond the automatic license's term. If you need more time, contact your Pentaho sales representative or pre-sales support engineer.

Cannot Create Hibernate Tables in MySQL

The Pentaho solution repository uses long text strings that require a longer maximum character limit than the default UTF-8 configuration allows. Therefore, if your MySQL character set is configured to use UTF-8, you must change it

to **ASCII** instead in order to use it as a Pentaho solution database. Using UTF-8 will prevent the MySQL initialization scripts from running during installation.

Case-Insensitivity for Usernames, Passwords, and Filenames

This problem can manifest in two different known ways. It can show as case-insensitivity as mentioned above, or it can show up when you change the case of an action sequence file name.

The source of the problem is the collation method and character set for your server. Selecting the wrong configuration will cause all usernames, passwords, and content files stored in the database to become identical in terms of characters without regard for case; this can cause name collisions. **SUZY**, **Suzy**, and **suzy** will all be the same username as far as the BA Server will be able to tell. To change this, you must set your database or table collation method to **latin1_general_cs**. You can do this by editing the database scripts mentioned below, or by modifying your server configuration.

Unable to Use the Database Init Scripts for PostgreSQL

The **pg_hba.conf** file contains host-based authentication information. If you can't run the SQL scripts that generate the Hibernate and Quartz databases, it's probably because the default user accounts for each database don't have the right permissions. To change this, edit the file to ensure that connections from local users created by the Pentaho sql scripts (**hibuser** and **pentaho_user**) will be able to connect. The default on Debian-based systems is for local connections to use **ident** authentication, which means that database users must have local user accounts. In other words, to continue using **ident**, you would have to create local **hibuser** and **pentaho_user** accounts. It's easier to just change the authentication method to something less restrictive, if your IT manager permits you to do so.

HTTP 500 or "Unable to connect to BA Server" Errors When Trying to Access Enterprise Console

If you have trouble controlling the BA Server through the Pentaho Enterprise Console, it is likely because you've changed the BA Server's IP address or hostname (the **fully-qualified-server-url** node in **web.xml**) from the default 127.0.0.1 to something else without also applying this change to the Proxy Trusting Filter in the **web.xml** file. See [Configuring the Proxy Trusting Filter](#) on page 13 for instructions on how to fix this.

Configuring the Proxy Trusting Filter

If you have set the BA Server to run on a specific IP address or hostname other than the default 127.0.0.1, you must modify the trusted proxy between the BA Server and Pentaho Enterprise Console to match that address or hostname.

1. Stop the Pentaho Enterprise Console.
2. Stop the BA Server.
3. Open **biserver-ee/tomcat/webapps/pentaho/WEB-INF/web.xml** and search for **TrustedIpAddr**s.



Note: The param-value immediately below **TrustedIpAddr**s is a comma-separated list of IP addresses that should be trusted.

```
<filter>
  <filter-name>Proxy Trusting Filter</filter-name>
  <filter-class>org.pentaho.platform.web.http.filters.ProxyTrustingFilter</filter-
class>
  <init-param>
    <param-name>TrustedIpAddr</param-name>
    <param-value>127.0.0.1</param-value>
    <description>Comma separated list of IP addresses of a trusted hosts.</
description>
  </init-param>
</filter>
```

4. Add the IP address of the host running the Pentaho Enterprise Console.
5. Open **enterprise-console/resource/config/console.xml** and search for **platform-username**.

6. Change the default **admin** user name, (default is joe) to your desired admin user name.
7. Start the BA Server.
8. Start the Pentaho Enterprise Console.

JBoss Fails to Start When the Pentaho HSQLDB Sample Database Is Running



Note: This problem can also manifest as the Pentaho sample database refusing to start when the BA Server is deployed to JBoss.

The Pentaho-supplied HSQLDB sample database operates on the default HSQLDB port of 9001. JBoss has its own HSQLDB instance running on the same port; therefore, the port collision will prevent the JBoss version from starting, and cause the startup process to halt. You can change the Pentaho sample database port by editing the **start_hypersonic** script and adding the **-port 9002** switch to the last line:

```
"$_PENTAHO_JAVA" -cp $THE_CLASSPATH org.hsqldb.Server -port 9002 -database.0 $DIR_REL/
hsqldb/sampledatab -dbname.0 sampledatab -database.1 $DIR_REL/hsqldb/hibernate -dbname.1
hibernate -database.2 $DIR_REL/hsqldb/quartz -dbname.2 quartz
```

JBoss Fails to Start After Manually Unpacking pentaho.war

If you unpack the pentaho.war file by hand, you must name the resultant directory **pentaho.war** as well. If you unpack it to any other directory name, including "pentaho" without the .war extension, JBoss will fail to deploy the WAR without any meaningful warnings.

Web-App Path doesn't validate in Enterprise Console

If you are deploying to JBoss and don't unpack the pentaho.war into a directory, the Pentaho Enterprise Console will not be able to find the correct Web-App Path value. To fix the problem, unpack the pentaho.war to a directory of the same name.

LDAP incorrectly authenticates user IDs that don't match letter case

Some LDAP implementations are case-insensitive, most notably Microsoft Active Directory. When using one of these LDAP distributions as a BA Server authentication backend, you might run into an issue where a valid username with invalid letter cases will improperly validate. For instance, if **Bill** is the valid user ID, and someone types in **bill** at the Pentaho User Console login screen, that name will authenticate, but it might have improper access to parts of the BA Server.

The fix for this is documented: [Forcing Case-Sensitivity in LDAP](#) on page 14.

Forcing Case-Sensitivity in LDAP

Some LDAP implementations (such as Active Directory) are case-insensitive regarding user names. When using one of these LDAP distributions for your Pentaho authentication backend, when the username typed into the Pentaho User Console login screen doesn't exactly match the case of the user ID in the directory, the directory server will positively authenticate it. However, the user ID that gets passed to the BI Platform contains the exact username that was typed, not the one in the directory. This is a potential security risk. In order to force case sensitivity, follow the below instructions.

1. Stop the BA Server.
2. Edit the `/pentaho/server/biserver-ee/pentaho-solutions/system/applicationContext-spring-security-ldap.xml` file.
3. Find the **daoAuthenticationProvider** bean, and below the last **</constructor-arg>** therein, and add the **<property>** definition shown in the example:

```
<property name="userDetailsContextMapper">
  <ref local="ldapContextMapper" />
</property>
```

```
</property>
```

4. After the `</bean>` tag for **daoAuthenticationProvider**, add the following bean definition, changing the **ldapUsernameAttribute** from **samAccountName** to the value that matches your environment:

```
<bean id="ldapContextMapper"
  class="org.pentaho.platform.engine.security.UseridAttributeLdapContextMapper">
  <property name="ldapUsernameAttribute" value="samAccountName" />
</bean>
```

5. Start the BA Server.

The BA Server will now force case sensitivity in LDAP usernames.

LDAP Roles Are Not "Admin" and "Authenticated"

You must not use **Admin** and **Authenticated** roles in your LDAP. Instead you must configure your system to use **pentahoAdmins** and **pentahoUsers** or other easily identifiable role names. Edit `/pentaho-solutions/system/applicationContext-spring-security.xml`. At the bottom of this file, you will find a number of lines that look like: `A/docs/. *Z=Anonymous,Authenticated`.

These are entries for URL Security. They are regular expressions to match a path on the browser's URL that require the user to be a member of the defined role to gain access. In the example above, both **Anonymous** and **Authenticated** get access. In the example above, use **pentahoUsers** in the place of **Authenticated**. by entering `A/docs/. *Z=Anonymous,pentahoUsers`. For all entries that show **Authenticated**, replace it with **pentahoUsers** or your chosen name. Replace **Admin** with **pentahoAdmins** or your chosen name. For the change from **Authenticated** to **pentahoUsers** replace all occurrences. For **Admin** to **pentahoAdmins** you need to be a little more careful because there are some entries that look like this: `A/admin. *Z=pentahoAdmins`.

Edit the `/pentaho-solutions/system/repository.spring.xml` file and change:

```
<bean id="singleTenantAuthenticatedAuthorityName" class="java.lang.String">
  <constructor-arg value="Authenticated" />
</bean>
```

to:

```
<bean id="singleTenantAuthenticatedAuthorityName" class="java.lang.String">
  <constructor-arg value="pentahoUsers" />
</bean>
```

and:

```
<bean id="singleTenantAdminAuthorityName" class="java.lang.String">
  <constructor-arg value="Admin" />
</bean>
```

to:

```
<bean id="singleTenantAdminAuthorityName" class="java.lang.String">
  <constructor-arg value="pentahoAdmins" />
</bean>
```

With LDAP Authentication, the PDI Repository Explorer is Empty

If you log into an enterprise repository from Spoon before you switch the authentication backend to LDAP, then the repository IDs and security structures will be broken. You won't see an error message, but the enterprise repository explorer will be empty and you won't be able to create new folders or save PDI content to it. To fix the problem, you will have to delete the security settings established with the previously used authentication method, which will force the DI Server to regenerate them for the LDAP backend.



Danger: Following this procedure will destroy any previously defined enterprise repository users, roles, and access controls. You should back up the files that you're instructed to delete below.

1. Stop the DI Server
2. Delete the security and default directories from the following directory: `/pentaho-solutions/system/jackrabbit/repository/workspaces/`
3. Start the DI Server

You should now have a proper LDAP-based enterprise repository that can store content and create new directories.

Old Action Sequences Fail on Scrollable Result Sets



Note: This topic was previously titled: **Changing the Default Result Set Cursor Type.**

Most databases can generate a scrollable result set that can be traversed multiple times. The Pentaho BI Platform assumes that your data source returns scrollable result sets, but will fall back to a forward-only result set when it encounters an error condition that may be caused by requesting a scrollable result set from a data source that does not support that functionality. If the error persists, that means that there is a different problem, so the result set fails and the error is shown to the user.

Since version 2.0, the BI Platform has had the ability to request a forward-only result set if specifically told to do so; this would prevent a live dataset from being generated, even with a database that supports scrollable result sets. When the forward-only flag is absent in an action sequence, the default behavior as of BA Server 3.5 is to try a scrollable result set, and if there is an error, the BI Platform will switch to a forward-only result set. This means that action sequences (including published reports) created with the BA Server versions 2.0, 2.0.1, 3.0, or 3.0.1 may have failed in the past under the described circumstances, but would now succeed.

To change the default method of parsing a result set to the old way, edit the `/pentaho/server/biserver-ee/pentaho-solutions/system/pentahoObjectsSpring.xml` file, then set the **fallBackToNonscrollableOnError** to **false**.

Geo Maps partially rendered

If your Geo Map visualizations in Analyzer are not correctly displaying, it's possible that the BA Server is not giving them enough time to fully render. See [Setting the Geo Map Rendering Wait Time](#) on page 16 for more information.

Setting the Geo Map Rendering Wait Time

There is a static length of time that the BA Server waits for a Geo Map to finish rendering before the image is captured; this allows for all map tiles to be downloaded and points to be plotted. The default wait time is **1200 milliseconds**, and it is set through the **<map-export-javascript-delay>** node in the `/pentaho-solutions/system/pentaho-geo/settings.xml` file:

```
<map-export-javascript-delay>1200</map-export-javascript-delay>
```

Mac OSX Error When Running Apps

When you run a Pentaho App file on OSX Mountain Lion, you may get a warning message from Gatekeeper. Temporarily changing security settings enables you to open apps successfully. For more information, see [OS X: About Gatekeeper](#).

1. Go to **System Preferences > Security > Privacy > General > Allow applications downloaded from**.
2. Change the setting to **Anywhere**.
3. Double-click the app.
A message box appears indicating this app is not signed.
4. Click **Open** to run the app.
5. Reset the security settings to what they originally were.

The next time you run this app, the message box does not appear and the app runs normally.

Report Designer and Reporting Engine

This section contains problems and solutions that pertain to Pentaho Report Designer.

Enabling Multi-Valued report Parameters for Metadata-based Queries Created with Previous Versions of Report Designer

In versions 3.7 and prior, there was no support for multi-value parameters in a Metadata query. If you have a report created in an earlier version, which contains a Metadata query and an "exactly matches" condition, the report will continue to work as is; however, if you try to change the parameter from a drop-down to a multi-selection type, such as a checkbox containing *more than one value*, the report will fail.

To resolve the problem, simply open the query for editing (Query Editor) and click **OK**. This adjusts MQL query to use the EQUALS function instead of the = operator. No additional changes are necessary.

Report Elements With Dynamic Heights Overlap Other Elements

If you have overlapping elements in your report whenever you use the **dynamic-height** style property, or if you'd like to create proper table rows so that elements of the second row get pushed down by the expanding content of the first row, then follow the directions below to create a two-row details band.

1. Select your **Details** band in the **Structure** pane, then go to the **Style** pane and change the value of **layout** to **block**.
2. Right-click the **Details** band, then select **band** from the **Add Element** context menu.
3. Move or add the elements for the first row into the band you just created.
4. Add another band, then move or copy all elements for the second row into the second band.

When your first row elements expand, your second row elements will be pushed down. Repeat this process as necessary for multiple rows.

Null Pointer Exception When Sharing Result Sets

A null pointer exception error may occur in the Pentaho User Console when you share a relational result set (SQLLookupRule used multiple times) among inputs in a Secure Filter component. The error may be associated with Result Set Type in Pentaho Design Studio. If the **Live Result Set**, is enabled, (this is default behavior), the result set is available one time only; the data for the query is fetched once and cannot be reused. To solve the issue, enable **In-Memory Result Set**; this stores the data in memory so that it can be reused and prevents the error from occurring.

Columns Unexpectedly Merge When Exporting to Excel

If you export content from Report Designer to Excel, and end up with unexpectedly merged columns in the output, there is probably a horizontal alignment problem with your column header or footer labels. If a label spans two columns -- even by a tiny amount -- then the Reporting engine will force the two columns to merge in the output.

Check your horizontal elements for column overlap. If you need more information on this topic, refer to the **Aligning Elements** section of the *Pentaho Report Designer User Guide*.

Analysis

This section contains problems and solutions that pertain to Pentaho Analysis, including JPivot, Schema Workbench, and Pentaho Analyzer.

Old Analysis Schemas Still Show Up in Pentaho User Console

If you have unwanted analysis schemas in the schema list in the Pentaho User Console -- this is the list that appears when you create a new analysis view or Analyzer report -- then you must edit the `datasources.xml` file and remove the unwanted entries as explained in [Removing Mondrian Data Sources](#) on page 19.

Removing Mondrian Data Sources

Every time you publish a schema from Schema Workbench or Pentaho Data Integration, a new XMLA data source entry is created in `/pentaho/server/biserver-ee/pentaho-solutions/system/olap/datasources.xml`. If you modify a schema and republish it, the data source entry will be updated accordingly. You do not have to manually add anything to this file, and under most circumstances you won't have to modify or remove anything from it, either. However, there is no automatic method to expire a registered data source, so each time you publish a new schema, a new entry is permanently added to `datasources.xml`.

The Pentaho User Console uses `datasources.xml` to populate a schema drop-down selection list that appears when users create new analysis views and Analyzer reports. As you phase out old analysis schemas, you will have to manually remove their `<DataSource>` entries in `datasources.xml`.

Below is the example `datasources.xml` that ships with a default Pentaho configuration:

`datasources.xml`

```
<?xml version="1.0" encoding="UTF-8"?>
<DataSources>
  <DataSource>
    <DataSourceName>Provider=Mondrian;DataSource=Pentaho</DataSourceName>
    <DataSourceDescription>Pentaho BI Platform Datasources</DataSourceDescription>
    <URL>http://localhost:8080/pentaho/Xmla?userid=joe&password=password</URL>
    <DataSourceInfo>Provider=mondrian</DataSourceInfo>
    <ProviderName>PentahoXMLA</ProviderName>
    <ProviderType>MDP</ProviderType>
    <AuthenticationMode>Unauthenticated</AuthenticationMode>
    <Catalogs>
      <Catalog name="SteelWheels">
        <DataSourceInfo>Provider=mondrian;DataSource=SampleData</DataSourceInfo>
        <Definition>solution:steel-wheels/analysis/
steelwheels.mondrian.xml</Definition>
      </Catalog>
      <Catalog name="SampleData">
        <DataSourceInfo>Provider=mondrian;DataSource=SampleData</DataSourceInfo>
        <Definition>solution:steel-wheels/analysis/SampleData.mondrian.xml</
Definition>
      </Catalog>
    </Catalogs>
  </DataSource>
</DataSources>
```

Multi-Byte Characters Don't Appear In PDFs Exported From Analyzer

If you are using a multi-byte character set, such as would be used for languages like Japanese or Chinese, and find that you have missing or corrupted output when exporting Analyzer reports to PDF, you will have to specify a default TrueType font for PDF rendering that supports multi-byte characters. The default PDF font in Analyzer is Helvetica, which does not support multi-byte character sets. To change this setting, see [Setting a Default Font for PDF Exports](#) on page 20.

Setting a Default Font for PDF Exports

Before following this procedure, stop the BA Server and Pentaho Enterprise Console.

When displaying data in Analyzer, your reports will use the default browser fonts. However, the PDF export function may not have the same fonts available to it when creating a PDF from your Analyzer report, resulting in output that doesn't look the same way in PDF format as it does in the browser. The default font for PDFs is Helvetica, but you can specify any TrueType font or collection to replace it. Follow the instructions below to specify a different font for PDF exports.



Note: If you have localized your schema in a language that uses a multi-byte character set (most Asian languages fit into this category), this process is required to make PDF output appear without errors.

1. Edit the **analyzer.properties** file in the `/pentaho/server/biserver-ee/pentaho-solutions/system/analyzer/` directory.
2. Uncomment the **renderer.pdf.font.path** line.

```
renderer.pdf.font.path=C:/WINDOWS/Fonts/MSGOTHIC.TTC,1
```

3. Replace the value of this line with the TrueType font or collection that you want to use as the default.
If you are specifying a collection, you must put a `,1` after the font name, as shown in the above example. This does not apply to individual fonts (TTF files).

```
renderer.pdf.font.path=/usr/share/fonts/truetype/freefont/FreeSans.ttf
```

4. Save and close the file, and start the BA Server.

Your PDF exports from Analyzer should have the font you specified.

Data Integration

This section contains problems and solutions that pertain to Pentaho Data Integration.

Troubleshooting a Corrupted DI Server Repository

If the DI Server's enterprise repository becomes corrupt, it will be unresponsive, content may be missing or inaccessible, and an error message similar to this will appear in the `/data-integration-server/tomcat/logs/catalina.out` log file:

```
ERROR [ConnectionRecoveryManager] could not execute statement, reason: File corrupted
while reading record: "page[48970] data leaf table:8 entries:1 parent:49157 keys:
[118547] offsets:[737]". Possible solution: use the recovery tool [90030-131], state/
code: 90030/90030
```

If this happens, shut down the DI Server and restore your enterprise repository from a recent backup.


If you do not have a viable backup, you may be able to minimize data loss by identifying the exact file that is corrupt. To do this, enable debug logging by adding the following XML snippet above the `<root>` element in the `/WEB-INF/classes/log4j.xml` inside your deployed `pentaho.war`:

```
<category name="org.pentaho.platform">
  <priority value="DEBUG"/>
</category>
```

Restart the DI Server and retry the action that caused the original error. If it occurs again, shut down the DI Server and open the **catalina.out** log file in Tomcat. Look for the last line that appears before the error; it usually contains the name of the file that has been damaged. When you are finished investigating the data corruption, remove the extra logging capabilities so that your DI Server log files don't become large and unmanageable.


```
reading file with id 'xyz' and path '/public/a.txt'
```

You can also try to recover your PDI data from the damaged database by using a recovery program, as explained in [Using the H2 Database Recovery Tool](#) on page 21.

 **Note:** If the database has become corrupt, the damaged rows will not be exported with the recovery tool, and whatever information they contained will be lost.

Using the H2 Database Recovery Tool

The DI Server includes a third-party H2 database recovery tool that enables you to extract raw data from your PDI enterprise repository. This is primarily useful in situations where the repository has become corrupt, and you don't have any relevant backups to restore from.

 **Note:** If the database has become corrupt, the corrupted rows will not be exported. Any information contained in corrupted rows is unrecoverable through this method.

The recovery tool is a JAR that must be run via the Java command. The output is a SQL dump that you can then attempt to import after you've re-initialized your DI Server database.

You can read more about the recovery tool on the H2 Web site: http://www.h2database.com/html/advanced.html#using_recover_tool.

Follow the directions below to use the recovery tool.

1. Open a terminal on (or establish an SSH connection to) your DI Server.
2. Navigate to the `/pentaho-solutions/system/jackrabbit/repository/version/` directory.

```
cd data-integration-server/pentaho-solutions/system/jackrabbit/repository/version/
```

3. Run the **h2-1.2.131.jar** H2 database JAR with the recovery tool option.

```
java -cp h2-1.2.131.jar org.h2.tools.Recover
```

4. Create a directory to move your old database files to.

```
mkdir old
```

5. Move the old database files to the directory you just created.

```
mv db.h2.db db.trace.db old
```

6. Re-initialize the repository with the **RunScript** option, using the salvaged SQL dump as the source.

```
java -cp h2-1.2.131.jar org.h2.tools.RunScript -url jdbc:h2:./db -user sa -script db.h2.sql
```

7. The backup directory you created earlier (**old** in the above example) can be removed after you're sure that you don't need the corrupted database files anymore. However, it's a good idea to keep this around just in case you need it later.

You've successfully extracted all of the usable data from your corrupt enterprise repository, removed the damaged database files, and re-initialized the repository with the salvaged data.

Start the DI Server and check for further errors. If repository errors persist, contact Pentaho support and request developer assistance.

Action Sequences That Call PDI Content Won't Run

If you've established an enterprise repository in PDI to store your jobs and transformations, and you attempt to use that stored PDI content in an action sequence on the BA Server, the action sequence will not execute. This is because the BA Server needs specific connection information for the Data Integration (DI) Server in order to retrieve the job or transformation.

Adding PDI Enterprise Repository Content Support to the BA Server

If you are using a Pentaho Data Integration (PDI) enterprise repository (through a Data Integration Server) to store PDI jobs and transformations, and you plan on using those jobs and transformations in action sequences that will be run on the BA Server, you must install some BA Server plugins from the PDI client tool package. This is not a typical scenario, but there is no harm in performing it if you aren't sure of the details.

1. Download a PDI Enterprise Edition 4.4 client tool archive package from the Pentaho Customer Support Portal.
The package name (available in both tar.gz and zip formats) is: **pdi-ee-client-4.4.0-GA**
2. Unpack the archive to a temporary location.
3. Edit the `/pentaho/server/biserver-ee/pentaho-solutions/system/kettle/settings.xml` file.
4. Change the value of the `<repository.type>` node from **files** to **rdbms**.
5. Enter your enterprise repository connection information in the proper nodes.
6. Enter the location of your local **repositories.xml** file in the `<repositories.xml.file>` node.



Note: This file is created on your PDI client workstation when you establish a connection to an enterprise repository. Once you have made all of your repository connections on a workstation, copy the **repositories.xml** file to the `~/ .kettle/` directory on the BA Server and DI Server machines. If the client tool and servers are all on the same machine, you do not have to copy the file. If you have not yet established any repositories, you will have to revisit this procedure later when your PDI environment is fully configured.

7. Copy the contents of `/data-integration/plugins/` to the `/pentaho/server/biserver-ee/pentaho-solutions/system/kettle/plugins/` directory.

```
cp -r /tmp/data-integration/plugins/* /home/pentaho/pentaho/server/biserver-ee/pentaho-solutions/system/kettle/plugins/
```

8. Remove the unpacked archive.

```
rm -rf /tmp/data-integration/
```

Your BA Server is now configured to run content stored in the DI Server.

Jobs scheduled on the DI Server cannot execute a transformation on a remote Carte server

You may see an error like this one when trying to schedule a job to run on a remote Carte server:

```
ERROR 11-05 09:33:06,031 - !
UserRoleListDelegate.ERROR_0001_UNABLE_TO_INITIALIZE_USER_ROLE_LIST_WEBSVC!
    com.sun.xml.ws.client.ClientTransportException: The server sent HTTP
    status code 401: Unauthorized
```

To fix this, follow the instructions in [Executing Scheduled Jobs on a Remote Carte Server](#) on page 23

Executing Scheduled Jobs on a Remote Carte Server

Follow the instructions below if you need to schedule a job to run on a remote Carte server. Without making these configuration changes, you will be unable to remotely execute scheduled jobs.



Note: This process is also required for using the DI Server as a load balancer in a dynamic Carte cluster.

1. Stop the DI Server and remote Carte server.
2. Copy the **repositories.xml** file from the `.kettle` directory on your workstation to the same location on your Carte slave.

Without this file, the Carte slave will be unable to connect to the enterprise repository to retrieve PDI content.

3. Open the `/pentaho/server/data-integration-server/tomcat/webapps/pentaho-di/WEB-INF/web.xml` file with a text editor.
4. Find the **Proxy Trusting Filter** filter section, and add your Carte server's IP address to the **param-value** element.

```
<filter>
  <filter-name>Proxy Trusting Filter</filter-name>
  <filter-class>org.pentaho.platform.web.http.filters.ProxyTrustingFilter</filter-
class>
  <init-param>
    <param-name>TrustedIpAddr</param-name>
    <param-value>127.0.0.1,192.168.0.1</param-value>
    <description>Comma separated list of IP addresses of a trusted hosts.</
description>
  </init-param>
  <init-param>
    <param-name>NewSessionPerRequest</param-name>
    <param-value>true</param-value>
    <description>true to never re-use an existing IPentahoSession in the
HTTP session; needs to be true to work around code put in for BISERVER-2639</
description>
  </init-param>
</filter>
```

5. Uncomment the proxy trusting filter-mappings between the `<!-- begin trust -->` and `<!-- end trust -->` markers.

```
<!-- begin trust -->
<filter-mapping>
  <filter-name>Proxy Trusting Filter</filter-name>
  <url-pattern>/webservices/authorizationPolicy</url-pattern>
</filter-mapping>

<filter-mapping>
  <filter-name>Proxy Trusting Filter</filter-name>
  <url-pattern>/webservices/roleBindingDao</url-pattern>
```

```

</filter-mapping>

<filter-mapping>
  <filter-name>Proxy Trusting Filter</filter-name>
  <url-pattern>/webservices/userRoleListService</url-pattern>
</filter-mapping>

<filter-mapping>
  <filter-name>Proxy Trusting Filter</filter-name>
  <url-pattern>/webservices/unifiedRepository</url-pattern>
</filter-mapping>

<filter-mapping>
  <filter-name>Proxy Trusting Filter</filter-name>
  <url-pattern>/webservices/userRoleService</url-pattern>
</filter-mapping>

<filter-mapping>
  <filter-name>Proxy Trusting Filter</filter-name>
  <url-pattern>/webservices/Scheduler</url-pattern>
</filter-mapping>

<filter-mapping>
  <filter-name>Proxy Trusting Filter</filter-name>
  <url-pattern>/webservices/repositorySync</url-pattern>
</filter-mapping>
<!-- end trust -->

```

6. Save and close the file, then edit the **carte.sh** or **Carte.bat** startup script on the machine that runs your Carte server.
7. Add **-Dpentaho.repository.client.attemptTrust=true** to the **java** line at the bottom of the file.

```

java $OPT -Dpentaho.repository.client.attemptTrust=true org.pentaho.di.www.Carte
"${1+$@}"

```

8. Save and close the file.
9. Start your Carte and DI Server

You can now schedule a job to run on a remote Carte instance.

Kitchen can't read KJBs from a Zip export



Note: This also applies to Pan and KTR files in Zip archives.

If you are trying to read a KJB file from a Zip export but are getting errors, you may have a syntax error in your Kitchen command. Zip files must be prefaced by a **!** (exclamation mark) character. On Linux and other Unix-like operating systems, you must escape the exclamation mark with a backslash: **\!**

Windows:

```

Kitchen.bat /file:"zip:file:///C:/Pentaho/PDI Examples/Sandbox/
linked_executable_job_and_transform.zip!Hourly_Stats_Job_Unix.kjb"

```

Linux:

```

./kitchen.sh -file:"zip:file:///home/user/pentaho/pdi-ee/my_package/
linked_executable_job_and_transform.zip\!Hourly_Stats_Job_Unix.kjb"

```


PDI Transformation Logging Doesn't Show In PEC

If you've enabled logging for a PDI transformation, then check the Pentaho Enterprise Console for log information, you may not find it if you didn't select all available fields, including **LOG_FIELD**, when you chose transformation fields.

To fix this, reconfigure the transformation logging options to include all available fields.

DI Server fails to start (IOException: Premature end of file)

If the DI Server fails to start, and you see an exception error like the one printed in the example, then follow the below directions to fix the **custom_nodetypes.xml** file.

1. Stop the DI Server.
2. Create a backup copy of the `/data-integration-server/pentaho-solutions/system/jackrabbit/` directory.
3. Delete the `/data-integration-server/pentaho-solutions/system/jackrabbit/` directory.
4. Create a new empty `/data-integration-server/pentaho-solutions/system/jackrabbit/` directory.
5. Create a new empty `/data-integration-server/pentaho-solutions/system/jackrabbit/repository/` directory.
6. Copy the **repository.xml** file from the directory you just backed up to the new jackrabbit directory.

```
cp ../backup_jackrabbit/repository.xml ./
```

7. Start the DI Server
8. Wait until the DI Server has fully started and initialized.
9. Stop the DI Server.
10. Copy the `/data-integration-server/pentaho-solutions/system/jackrabbit/repository/nodetypes/custom_nodetypes.xml` file to the same location in the backup directory you created earlier.

```
cp custom_nodetypes.xml ../../../../backup_jackrabbit/repository/repository/nodetypes/
```

11. Delete the newly-created `/data-integration-server/pentaho-solutions/system/jackrabbit/` directory.
12. Move the backup copy of the jackrabbit directory to its original location.

```
mv backup_jackrabbit /home/pentaho/pentaho/server/data-integration-server/pentaho-solutions/system/jackrabbit/
```

13. Start the DI Server

The DI Server should now start normally.

Using ODBC


Although ODBC can be used to connect to a JDBC compliant database, Pentaho does not recommend using it and it is not supported. For details, this article explains "Why you should avoid ODBC." <http://wiki.pentaho.com/pages/viewpage.action?pageId=14850644>.

Metadata

This section contains problems and solutions that pertain to Pentaho Metadata, including Metadata Editor and the Agile BI and Pentaho User Console components that auto-generate metadata models.

Managing Multiple Outer-Joins

When you have three or more tables that require outer joins, the *order* in which the tables are joined is critical. Consider the example below:



In the sample preview below, the entries, 1, 2 ,3, and 4, in Table4 are taken and outer-joined with the records in the two other tables. The three other tables contain fewer records. The relationships are defined but now the order of execution critical. Relationship A is executed first, followed by B, then C.

| Examine preview data | | | | | |
|--|-----------------|-----------------|-----------------|-----------------|--|
| Rows of step: The first 5000 rows from the generated query | | | | | |
| # | BC_TABLE1_T1_PK | BC_TABLE2_T2_PK | BC_TABLE3_T3_PK | BC_TABLE4_T4_PK | |
| 1 | 1 | 1 | 1 | 1 | |
| 2 | 2 | 2 | | 2 | |
| 3 | | | | 3 | |
| 4 | | | | 4 | |

Below is the query that is generated:

```

SELECT DISTINCT
    TABLE1.PRIMARYKEY AS COL0
    ,TABLE2.PRIMARYKEY AS COL1
    ,TABLE3.PRIMARYKEY AS COL2
    ,TABLE4.PRIMARYKEY AS COL3

FROM TABLE4 LEFT OUTER JOIN
    (
        TABLE3 FULL OUTER JOIN
        (
            TABLE1 FULL OUTER JOIN TABLE2
            ON ( TABLE1.PRIMARYKEY = TABLE2.FOREIGNKEY )
        )
        ON ( TABLE2.PRIMARYKEY = TABLE3.FOREIGNKEY )
    )
ON ( TABLE2.FOREIGNKEY = TABLE4.PRIMARYKEY )

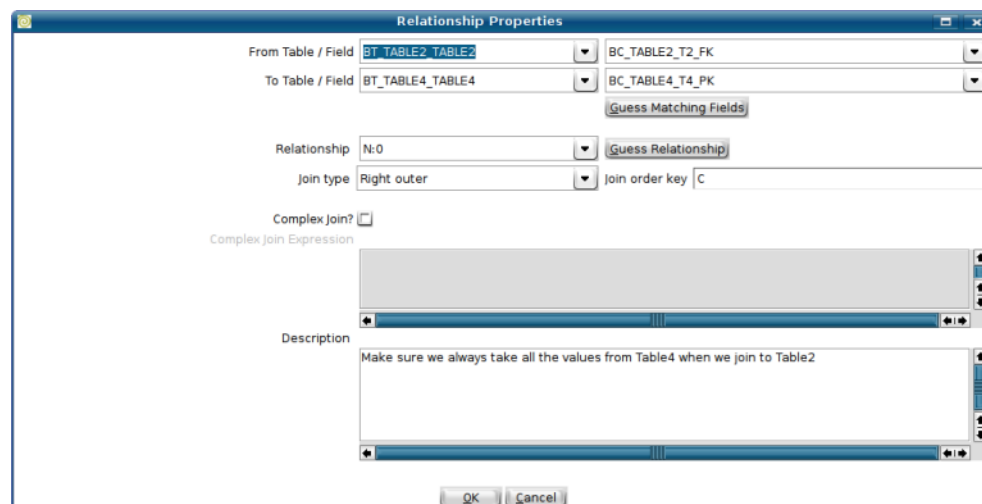
ORDER BY
    TABLE4.PRIMARYKEY

```

The nested join syntax that is generated forces the order of execution:

- Join Table1 and Table2 (shown in red)
- Join Table3 and A = B (shown in blue)
- Join Table4 with B = Result

Other orders of execution are just as valid depending on the business context to which they are applied. Another order of execution will generate a different result. To allow business model designers to ensure that user selections are executed in a specific way, a **Join Order Key** is added to the **Relationship Properties** dialog box.



The join order key is relevant only in instances in which outer joins are deployed in business models. To make the importance of the execution order apparent, this information is displayed in the graphical view of the model, as shown below.



Note: It is not mandatory to use uppercase letters, (A, B, C, as shown in the first image), to set the order in which tables are executed. Any alphanumeric characters (0-9, A-Z) can be used. The system will calculate the ASCII values of each character; the values are then used to determine the order of execution. In the example, A, B, C, AA, AB, Pentaho Metadata Editor will execute the table relationships in the following order: A, AA, AB, B, C.

Using the Delay Outer Join Conditions Property

To force conditions that would ordinarily be processed in the JOIN condition to be processed in the WHERE clause, follow the directions below to create a **delay_outer_join_conditions** custom property.

1. Right-click on a business model and select **Edit**.
2. Add a property by clicking the green **+** icon.
3. Select **Add a Custom Property** and set its ID to **delay_outer_join_conditions** and select **boolean** for the **Type**, then click **OK**.
4. Select the newly-created **delay_outer_join_conditions** property, then click the checkbox for **delay_outer_join_conditions** under the **Custom** heading on the right side of the window, then click **OK**.

Instead of the conditions being rolled into the JOIN clause, they will be allowed to roll down into the WHERE clause.